

GradCafe and similar anonymous self-submission websites have some clear limitations as data sources. Since users choose to submit their own information, the dataset is not really a random or fully representative sample of all graduate school applicants. People may be more likely to post if they received a very exciting acceptance, a frustrating rejection, or if they are already the type of applicant who spends time on admissions forums. This can create some selection bias in the data. There are also a lot of missing values because not every person reports their GPA, GRE scores, nationality, comments, or exact program details. Some entries may also be repeated, exaggerated, typed incorrectly, or formatted in slightly different ways. Since the submissions are anonymous, there is not a reliable way to verify that every entry is accurate.

Because of that, my analysis results should be treated as a summary of the GradCafe entries I scraped, not as official admissions statistics. For example, the average GRE scores in this dataset may be higher than national averages because the people posting on GradCafe are already a self selected group of graduate school applicants, and applicants with stronger scores may be more comfortable sharing them. It also makes sense that competitive programs and highly invested applicants might be overrepresented on the site. The results did not completely surprise me for that reason, since online admissions forums tend to attract people who are especially focused on the application process. Compared with official university or testing agency data, GradCafe data is less standardized, less complete, and less reliable, but it can still be useful for noticing general patterns as long as those limitations are kept in mind.